

# CS 101: Computer Programming and Utilization

Jul-Nov 2017

Umesh Bellur  
([cs101@cse.iitb.ac.in](mailto:cs101@cse.iitb.ac.in))

Lecture 15: Standard Library

# About These Slides

- Based on Chapter 22 of the book  
*An Introduction to Programming Through C++*  
by Abhiram Ranade (Tata McGraw Hill, 2014)
- Original slides by Abhiram Ranade
  - First update by Uday Khedker
  - Second update by Sunita Sarawagi

# The Standard Library

- Comes with every C++ distribution
- Contains many functions and classes that you are likely to need in day to day programming
- The classes have been optimized and debugged thoroughly
- If you use them, you may be able to write programs with very little work
- Highly recommended that you use functions and classes from the standard library whenever possible

# Outline

- The string class
- The template class vector
  - Multidimensional vectors
  - Sorting a vector
- The template class map
  - Iterators
- Remarks

# The String Class

- A much more powerful version of the String class developed in Chapter 21
- More constructors
- Concatenation using +
- Works with >> and <<
- Operations for extracting substrings and finding one string inside another

# Examples

```
#include <string>    // Needed to use the string class
string v = "abcdab"; // constructor
string w(v);         // another constructor. w = v
v[2] = v[3];         // indexing allowed. v becomes "abddab"
cout << v.substr(2) << v.substr(1,3) << endl;
    // substring starting at v[2] ("ddab")
    // Substring starting at v[1] of length 3 ("bdd")
int i = v.find("ab"); // find occurrence of "ab" in v
                    // and return index
int j = v.find("ab",1); // find from index 1
cout << i << ", " << j << endl; // will print out 0, 4.
```

# Remarks

- If the find member function does not find the argument in the receiver, then it returns a constant `string::npos`, which is a value which cannot be a valid index
  - You can determine whether the argument was found by checking whether the returned index equals `string::npos`
- `string s; s.size() or s.length()` returns a value of type `size_t` (a redefinition of unsigned int)
- A string object can be passed by value, in which case it is copied, or by reference
- More details on the web. Example:  
<http://www.cplusplus.com/reference/string/>

# Templates

- Function templates (Sec 12.5 in book)
- Consider these three functions: same body, different types

```
int Abs(int x) {  
    if (x < 0)  
        return -x;  
    else return x;  
}
```

```
float Abs(float x) {  
    if (x < 0)  
        return -x;  
    else return x;  
}
```

```
double Abs(double x) {  
    if (x < 0)  
        return -x;  
    else return x;  
}
```

A common template to unite them all....

```
template<typename T>  
T Abs(T x) {  
    if (x < 0)  
        return -x;  
    else return x;  
}
```



# Template Class

- Like function templates, create class with templates.

```
template <class T>
class Queue {
    int front, nWaiting;
    T elements[100];
public:
    bool insert(T value)
    {...}
    bool remove(T &val)
    {...}
};
```

```
main () {
    Queue<V3> q;
    Queue<int> r;
}
```

# The Template Class Vector

- Friendlier, more versatile version of arrays
- Must include header file `<vector>` to use it
- You can make vectors of any type by supplying the type as an argument to the template
- Indexing possible like arrays
- Possible to extend length, or even insert in the middle
- We will not discuss how the vector class is implemented, but you should be able to guess that its member functions would allocate memory and deallocate it as needed

# Examples

```
#include <vector> // needed
vector<int> v1;    //empty vector. Elements will be int
vector<float> v2; //empty vector. Elements will be float
vector<short> v3(10); // vector of length 10.
                    // Elements are of type short
vector<char> v4(5,'a'); // 5 elements, all 'a'
cout << v3.size() << endl; // prints vector length, 10
                    // v3.length() is same
v3[6] = 34;         // standard indexing
```

## Examples (Contd.)

```
#include <vector>           // needed
...

v3.push_back(22);           // append 22 to v3.
                             // Length increases

vector<char> w;
w = v5;                     // element by element copy
v1.resize(9);               // change length to 9
v2.resize(5, 3.3);          // length becomes 5, all
                             // values become 3.3

vector<string> s;           // vector of string
vector<vector<int> > vv;    // allowed!
```

# A Technical Remark

- The member function `size` returns a value of type `size_t`
- `size_t` is an unsigned integer type; it is meant specially for storing array indices
- When going through array elements, use `size_t` for the index variable

```
vector<double> v(10);           // initialize v
for(size_t i=0; i<v.size(); i++)
    cout << v[i] << endl;
```

- If `i` were declared `int`, then the compiler would warn about the comparison between `i` and `v.size()`
  - comparison between signed and unsigned int, which is tricky as discussed in Section 6.8.
  - By declaring `i` to be `size_t`, the warning is suppressed.

# Multidimensional Vectors

```
vector<vector <int> > vv;  
// each element of vv is itself a vector of int  
// we must supply two indices to get to int  
// Hence it is a 2d vector!  
// Currently vv is empty  
vector<vector <int> > vv1(5, vector<int>(10,23));  
// vv1 has 5 elements  
// each of which is a vector<int>  
// of length 10,  
// having initial value 23
```

# Multidimensional Vectors

- Note that the syntax is not new/special
- It is merely repeated use of specifying the length and initial value:
- `vector<type> name(length, value)`
- Two dimensional arrays can be accessed by supplying two indices, i.e. we may write `vv1[4][6]` and so on
- Write `vv1.size()` and `vv1[0].size()` to get number of rows and columns

# Creating A 5x5 Identity Matrix

```
vector<vector<double>> m(5, vector<double>(5,0));
```

```
    // m = 5x5 matrix of 0s
```

```
    // elements of m can be accessed
```

```
    // by specifying two indices
```

```
for(int i=0; i<5; i++)
```

```
    m[i][i] = 1;
```

```
    // place 1s along the diagonal
```



# Remarks

- The book gives a matrix class which internally uses vector of vectors
- This class is better than two dimensional arrays because it can be passed to functions by value or by reference, with the matrix size being arbitrary

# Sorting A Vector

- C++ provides a built-in facility to sort vectors and also arrays
- You must include `<algorithm>` to use this

```
vector<int> v(10);  
// somehow initialize v  
sort(v.begin(), v.end());
```

- That's it! v is sorted in non decreasing order
- `begin` and `end` are “**iterators**” over v. Think of them as abstract pointers to the beginning and the end.

# Sorting An Array

- The algorithms in header file `<algorithm>` can also sort arrays as follows

```
double a[100];  
// somehow initialize a  
sort(a, a+100); // sorted!  
// second argument is name+length
```

- More variations in the book

# The Map Template Class

- A vector or an array give us an element when we supply an index
  - Index must be an integer
- But sometimes we may want to use indices which are not integers, but strings
  - Given the name of a country, we may want to find out its population, or its capital
  - This can be done using a `map`

# Map: General Form And Examples

- General form:

`map<indextype, valuetype> mapname;`

- Examples:

`map<string,double> population;`

Indices will have type `string` (country names), and elements will have type `double` (population)

`map<string, vector<string>> dictionary;`

Maps words to their meanings expressed as a vector of other words.

# Using A Map

```
map<string,double> population;

population["India"] = 1.21;
                // in billions. Map entry created
population["China"] = 1.35;
population["USA"] = 0.31;

cout << population["China"] << endl;
                // will print 1.35

population["India"] = 1.22;
                //update allowed
```

# Checking if An Index is Defined

```
string country;  
cout << "Give country name: ";  
cin >> country;  
  
if(population.count(country)>0)  
    // true if element with index = country  
    // was stored earlier  
    // count is a known member function  
    cout << population[country] << endl;  
else cout << "Not known.\n";
```

# Remarks

- A lot goes on behind the scenes to implement a map
- Basic idea is discussed in Chapter 24 of our book
- If you wish to print all entries stored in a map, you will need to use iterators, discussed next



# Iterators

- A map can be thought of as holding a sequence of pairs, of the form (index, value)
- For example, the population map can be considered to be the sequence of pairs  
`[("China",1.35), ("India",1.21), ("USA", 0.31)]`
- You may wish to access all elements in the map, one after another, and do something with them
- For this, you can obtain an iterator, which points to (in an abstract sense) elements of the sequence

# Iterators

An iterator points to (in an abstract sense) elements of the sequence

- An iterator can be initialized to point to the **first** element of the sequence
- In general, given an iterator which points to some element, you can ask if there is any element **following** the element, and if so make the iterator point to the **next element**
- An iterator for a `map<index,value>` is an object with type `map<index,value>::iterator`

## Iterators (contd.)

- An iterator points to elements in the map; each element is a struct with members first and second
- We can get to the members by using dereferencing
- Note that this simply means that the dereferencing operators are defined for iterators
- If many elements are stored in an iterator, they are arranged in (lexicographically) increasing order of the key

# Example

```
map<string,double> population;  
population["India"] = 1.21;  
  
map<string,double>::iterator mi;  
mi = population.begin();  
    // population.begin() : constant iterator  
    // points to the first element of population  
    // mi points to (India,1.21)  
cout << mi->first << endl; // or (*mi).first << endl;  
    // will print out India  
cout << mi->second << endl;  
    // will print out 1.21
```

# Example

```
map<string,double> population;  
population["India"] = 1.21;  
population["China"] = 1.35;  
population["USA"] = 0.31;  
for(map<string,double>::iterator  
    mi = population.begin();  
    mi != population.end();  
        // population.end() : constant iterator  
        // marking the end of population  
    mi++)  
    // ++ just sets mi to point to the  
    // next element of the map  
    // loop body
```

## Example (Contd.)

```
map<string,double> population;
population["India"] = 1.21;
population["China"] = 1.35;
population["USA"] = 0.31;
for(map<string,double>::iterator
    mi = population.begin();
    mi != population.end();
    mi++)
{
    cout << (*mi).first << ": " << (*mi).second << endl;
    // or cout << mi->first << ": " << mi->second << endl;
}
// will print out countries and population in alphabetical order
```

# Remarks

- Iterators can work with vectors and arrays too
- Iterators can be used to find and delete elements from maps and vectors.

```
map<string,double>::iterator  
    mi = population.find("India");  
population.erase(mi);
```

# Map with user-defined class as index

- Any class used as indextype on a map must implement the "<" operator.
- Example, the following code will not work because "<" is not defined on V3.

- `class V3 {public: double x,y,z};`
- `map<V3, string> vec2string;`

- A correct implementation of V3 may be something like:

```
class V3 {
public:
double x,y,z;
bool operator<(const V3& a) const {
    if (x < a.x) return true;
    if (x == a.x && y < a.y) return true;
    if (x==a.x && y == a.y && z < a.z) return true;
    return false;
}
};
```



# Sets

- Sets are containers that store unique elements following a specific order
- The value of the elements in a set cannot be modified once in the container (the elements are always const), but they can be inserted or removed from the container
- Internally, the elements in a set are always sorted following a specific ordering criterion indicated by its internal comparison object

# Populating and Traversing a Set

```
#include <set>                // set class library

...

set<int> set1;                // create a set object,
                               // specifying its content as int
                               // the set is empty

int ar[]={3,2,4,2};
for (int i = 0; i < 4; i++) {
    set1.insert(ar[i]);    // add elements to the set.
}

for (set<int>::iterator iter = set1.begin(); iter != set1.end(); iter++) {
    cout << *iter << " ";
} // prints 2 3 4
```

# Application of Set

Given N students where each student has a list of courses that they have taken. Create group of all students that have taken exactly the same set of courses.

```
map<set<string>, vector<int>> study_group;
// key of the map is the set of courses.
// value is vector of student roll-numbers of students who
have taken this course.
cin >> N;
for(int i = 0; i < N; i++) {
    int roll, int n;
    cin >> roll >> n;
    set<string> subjects;
```

## Application of Set (continued)

```
for (int j = 0; j < n; j++) {  
    string s; cin >> s;  
    subjects.insert(s);  
}  
study_group[subjects].push_back(rollno);  
}
```

# List

- Implements a classic list data structure
- Supports a dynamic bidirectional linear list
- Unlike a C++ array, the objects the STL list contains cannot be accessed directly (i.e., by subscript)
- Is defined as a template class, meaning that it can be customized to hold objects of any type
- Responds like an unsorted list (ie. the order of the list is not maintained). However, there are functions available for sorting the list

# Populating and Traversing a List

```
#include <list>                // list class library
...
list<int> list1;                // create a list object,
                                // specifying its content as int
                                // the list is empty

for (i=0; i<5; i++)
    list1.push_back (i);        // add at the end of the list
...
while (list1.size() > 0)
{   cout << list1.front();      // print the front item
    list1.pop_front();          // discard the front item
}
```

# Concluding Remarks

- Standard Library contains other useful classes, e.g. queue, list, set etc.
- The Standard Library classes use heap memory, however this happens behind the scenes and you don't have to know about it
- The library classes are very useful. Get some practice with them
- More details on the web. Example:  
<http://www.cplusplus.com/reference/stl/>